Notes on Simulating Mutants: Efficient MSM-based Sampling Strategies

Robert T. McGibbon

I. INTRODUCTION

Mutational analysis is one of the central tools in experimental protein science, but remains a challenge for simulation. Using standard simulation approaches, the study of a protein and a single mutant requires 2x the computational effort of studying just the original protein. Can we use MSMs to do better?

Assume that we've run extensive, converged simulation of a protein, A, and we build an MSM. We now want to run a mutant, A'. We assume that A, and A' share a common state space – the interest is in how the mutation affects the transition probabilities.

If we assume that the mutation affects the transition matrix in a "small" way – that the mutation is appropriately classified as a perturbation, then it should be possible to "win". First, short trajectories (of length equal to a single lag-time) are sufficient, because (by assumption), we don't have to discover any new states. We just have to estimate the perturbed transition probabilities. Second, because the mutation is small, there's a lot of mutual information between transition probabilities in A and those in A'. We're not starting from scratch here.

II. PROBLEM STATEMENT

What is the mathematical model for the cross-talk between A and A'? Our observation of A must basically be the prior for A'.

Consider a row of the transition matrix $T^{A'}$, the transition probabilities leaving state i, $\vec{p}_i^{A'}$. In the absence of any simulation data on A', what is our prior distribution on $P(\mathbf{p_i}^{A'})$?

The simplest idea is that the $\vec{p}_i^{A'}$ are $\mathrm{Dir}(\alpha)$, where the α parameters (effectively the psuedocounts) are determined by the number of observed counts in protein A, $\vec{c}_i^A$. If there is a lot of mutual information between A and A', then the α should be equal to $\vec{c}_i^A$. This formalizes the idea that our predictions about protein A' made from data collected on A are equally confident as our predictions about protein A using the same data collected on A. But if there's not a lot of mutual information between the two proteins, then our prior on $\vec{p}_i^{A'}$ should be non-informative.

$$\vec{p}_i^{A'} \sim \mathrm{Dir}(q_i \cdot \vec{c}_i^A + 1/2)$$

Here, the parameter $q_i \in [0, 1]$ gives something like the expected strength of the information transfer between state i in the the two models. When $q_i = 1$, a count measured in A is "worth it's weight" in the A' model.

But as $q_i \to 0$, those counts are worthless in A' and we get a noninformative Jeffreys prior.

This statement of the problems is nice because it takes into account some uncertainty in the gold-standard model for A. It doesn't just look the MLE estimate of the transition matrix in A, but uses the counts directly to parameterize the distribution over A'. So for regions of the state space where $\vec{c}_i^{A'}$ is low, we're going to get a mostly uninformative prior naturally.

One question: should q be a single parameter for the whole model, or should every state have its own q_i? The later case could encode the idea that some of the states behave very similarly in A vs. A', whereas other might be very different.

Next: clearly q_i should not be a fixed parameter. It should also be a random variable, estimated from the data. Perhaps the appropriate prior on q_i is $q_i \sim \mathrm{Beta}(\alpha, \beta)$.

Now, if observe some outbound counts, $\vec{K}$, from state i in simulations of the mutant protein A' (note for consistency, we could notate $\vec{K}$ as $\vec{c}_i^{A'}$ instead), they are distributed as a multinomial with parameters $\vec{p}_i^{A'}$. The model specified is then:

$$\vec{c}_i^A = \text{prior observed counts in protein } A$$
$$\alpha, \beta = q\text{'s hyperparameters}$$
$$q_i \sim \mathrm{Beta}(\alpha, \beta)$$
$$\vec{p}_i^{A'} \sim \mathrm{Dirichlet}(q \cdot \vec{c}_i^A + 1/2)$$
$$K \sim \mathrm{Multinomial}(\vec{p}_i^{A'})$$

Can we do inference here? I think so. Let's start with the joint distribution of of $\vec{p}_i^{A'}$ and q_i. By bayes rule,

$$P(q_i, \vec{p}_i^{A'} \,|\, \vec{K}) \propto P(\vec{K} \,|\, q, \vec{p}_i^{A'})\, P(q_i, \vec{p}_i^{A'})$$

K is conditionally independent of q_i given $\vec{p}_i^{A'}$, and is multinomial. $P(q_i, \vec{p}_i^{A'})$ factors as $P(\vec{p}_i^{A'} \,|\, q_i) \cdot P(q_i)$, a Dirichlet times a Beta. Because of the conjugacy, $P(\vec{K} \,|\, q_i, \vec{p}_i^{A'})$ and $P(\vec{p}_i^{A'} \,|\, q_i)$ group together into an updated Dirichlet. Putting it all together, we get

$$P(q_i, \vec{p}_i^{A'} \,|\, \vec{K}) \propto \mathrm{Dir}(q_i \cdot \vec{c}_i^A + \vec{K} + 1/2) \cdot P_{\alpha, \beta}(q)$$

And then marginalizing out q_i, we have

$$P(\vec{p}_i^{A'} \,|\, \vec{K}) \propto \int_0^1 dq \; \mathrm{Dir}(q \cdot \vec{c}_i^A + \vec{K} + 1/2) \cdot P_{\alpha, \beta}(q_i)$$

III. ACTIVE LEARNING

We've done a little bit of simulation in A', and now we want to start some new simulations. From which state i should we start our next round of sampling?

The goal should be, in some sense, to maximize our expected information gain, or achieve the greatest reduction in the model's uncertainty. We consider two strategies alternative.

A. Maximize Information Gain in Transition Probabilities

First, we consider adaptively choosing the state from which to start new simulations in a way that maximizes the expected information gain in the transition probabilities p_{ij}. Conditional on q_i, the expected information gain in the outbound transition probabilities from state i is given by the expected Kullback–Leibler divergence between the current $P(p_i^{A'})$ and the updated posterior, $P(p_i^{A'}|e)$, where e is newly observed transition, distributed according to the current model.

The K-L divergence from one Dirichlet distribution, p, to another, q, is given by

$$D_{KL}(\lambda^q||\lambda^p) = \log\frac{\Gamma(\lambda^{qt})}{\Gamma(\lambda^{pt})} + \sum_{s=1}^{m} \log\frac{\Gamma(\lambda_s^p)}{\Gamma(\lambda_s^q)}$$
$$+ \sum_{s=1}^{m} [\lambda_s^q - \lambda_s^p] \left[\Psi(\lambda_s^q) - \Psi(\lambda^{qt})\right]$$

Where λ^p and λ^q are the parameters of the two Dirichlet distributions, $\lambda^{qt} = \sum_{s=1}^{m} \lambda_s^q$, $\lambda^{pt} = \sum_{s=1}^{m} \lambda_s^p$, and m is the number of states.

If λ^p is formed from λ^q by incrementing a single element l by one, then by simple properties of the gamma function, it can be shown that

$$D_{KL}(\lambda||\lambda + e_l) = -\log(\lambda^t) + \log(\lambda_l) + \Psi(\lambda^t) - \Psi(\lambda_l)$$

Then, the expected KL divergence is

$$E[D_{KL}] = \sum_{l} D_{KL}(\lambda||\lambda + e_l)P(l|\lambda)$$

$P(l\,|\,\lambda) = \int d\vec{p}\, P(l\,|\,\vec{p})P(\vec{p}\,|\,\lambda)$ follows the Multinomial-Dirichlet distribution, with $P(l\,|\,\lambda) = \lambda_l/\lambda^t$. Thus the expected information gain comes out to be

$$E[D_{KL}\,|\,q_i] = \Psi(\lambda^t) - \log(\lambda^t) + \frac{1}{\lambda^t} \sum_{l} [\lambda_l \left(\log(\lambda_l) - \Psi(\lambda_l)\right)]$$

This formula is still conditional on q_i. Thus the procedure is to fit the model for each row of the transition matrix with MCMC and then compute, for each sampled λ from the posterior, the expected information gain as above. The mean over the samples is then the final "score" for each state. At each step we choose the to initialize new sampling from the state with the highest such score.

B. Minimize Uncertainty in the First Eigenvalue

The procedure above treats the entropy in each row of the transition matrix independently and on equal footing. But if we care about predicting the long-time dynamics of the system accurately, we care substantially more about some states than others.

Instead of maximize the information gain in the transition probabilities, p_{ij}, an alternative would be to minimize the uncertainty in the model's longest implied timescale.